

UNIVERSITY OF PERADENIYA
Faculty of Engineering

END OF SEMESTER EXAMINATION, SEPTEMBER 2020

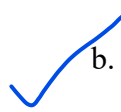
CO 527 ADVANCED DATABASE SYSTEMS
(Three Hours)

Registration Number:	E/...../.....
<div>1. Answer all questions on this booklet itself.</div> <div>2. All questions carry equal marks.</div>	

For Examiner's Use Only							
	Q1	Q2	Q3	Q4	Q5	Q6	Total
Marks Awarded							

1. Query Optimization and Database Tuning

-  a. What is a query tree? [10 Marks]

-  b. Compare the different approaches available for query optimization. [20 Marks]

- c. Consider the following relations and the SQL query based on the relations.

Relations:

EMPLOYEE (Ssn, Fname, Lname, Address, Bdate)

DEPARTMENT (Dnumber, Dname, Mgr_ssn)

PROJECT (Pnumber, Pname, Plocation, Dnum)

SQL query:

SELECT Pnumber, Dnum, Lname, Address, Bdate

FROM ((PROJECT JOIN DEPARTMENT ON Dnum=Dnumber)

JOIN EMPLOYEE ON Mgr_ssn=Ssn)

WHERE Plocation='Stafford';

- ✓ i. Draw an initial query tree to represent the query given above.

[10 Marks]

- ✓ ii. Show how the initial query tree is optimized based on the heuristics mentioned in (1.b) and find a final query tree that is efficient to execute. Briefly explain the steps.

[30 Marks]

- ✓ d. List three factors that influence physical database design. Briefly explain how those factors could be used to revise database design. [20 Marks]

- ✓ e. What are the reasons for tuning indexes when improving the overall throughput of transactions? [10 marks]

2. Transaction Processing

a. What is meant by a strict schedule?

[10 marks]

b. Is the following schedule with three transactions T_1 , T_2 and T_3 operating on three variables X, Y and Z a strict schedule? Explain.

$R_1(X), R_2(Z), R_1(Z), R_3(X), R_3(Y), W_1(X), c_1(\text{Commit}), W_3(Y),$
 $c_3(\text{Commit}), R_2(Y), W_2(Z), W_2(Y), c_2(\text{Commit});$

[30 Marks]

c. What is a serializable schedule?

[10 Marks]

d. Show that the following schedule with three transactions T_1 , T_2 and T_3 operating on three variables X, Y and Z is not serializable.

$R_2(Z), R_2(Y), W_2(Y), R_3(Y), R_3(Z), R_1(X), W_1(X), W_3(Y), W_3(Z), R_2(X),$
 $R_1(Y), W_1(Y), W_2(X)$

[30 Marks]

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and extend across the width of the page. There are no margins, text, or other markings on the paper.

e. Can the schedule in 2 (d) be made serializable? If so, how? [20 Marks]

[illegible]

3. Concurrency Control and Database Recovery Techniques

✓ a. How does two phase locking protocol guarantee serializability? [20 Marks]

[illegible]

- b. Show that the following schedule does not satisfy the two-phase locking protocol. [30 Marks]

T₁: Read_Lock (Y), T₁: Read_Item (Y), T₁: Unlock (Y), T₂: Read_Lock (X),
T₂: Read_Item (X), T₂: Unlock (X), T₂: Write_Lock (Y), T₂: Read_Item (Y),
T₂: Write_Item (Y), T₂: Unlock (Y)

- c. What is snapshot isolation? What are the advantages and disadvantages of concurrency control methods that are based on snapshot isolation? [20 Marks]

- d. Discuss the different types of transaction failures. What is meant by catastrophic failure? [20 Marks]

- e. Explain how to recover a database to most recent consistent state after a catastrophic failure. [10 Marks]

4. Distributed Database Systems

- a. Explain how fragmentation transparency will improve the performance of distributed database system. [10 Marks]

- b. Giving an example, explain horizontal fragmentation in distributed database systems. [20 Marks]

- c. What are the challenges in handling transactions in distributed database systems? [10 Marks]

- d. Explain how three-phase commit protocol handles the issues of two-phase commit protocol. [20 Marks]

- ii. Suppose that query Q is submitted at site 2. Identify the best strategy for executing the query. [10 Marks]

5. Big Data

- a. Define “big data”. [10 Marks]

- b. Briefly explain two example applications which use big data. [20 Marks]

c. Describe the characteristic *Veracity* in big data.

[20 Marks]

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

d. What is Apache Hadoop?

[25 Marks]

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

- e. Briefly explain what dimensions one should consider when deploying Hadoop in an organization? [25 Marks]

6. NoSQL databases

- a. For which types of applications were NoSQL systems developed? [20 Marks]

- b. What are the main categories of NoSQL systems? Provide an example for each category. [20 Marks]

-
- This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

- [illegible]

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

[20 Marks]

Blank lined paper for writing.

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

End of the Paper